

# Lock-Free Priority Queue Based on Multi-Dimensional Linked Lists

Jun Tao Luo

Yumin Chen

## Background

- Priority queues are central to many CS problems
- Investigate concurrent implementations
- Explore lock-free implementation using Multi-Dimensional Linked Lists (MDList)[1]

## Single Source Shortest Path

- Classic CS problem, compute the shortest paths from source node to all other nodes in a graph
- Sequentially solved using Dijkstra's algorithm
- Parallelized Dijkstra's algorithm [2]
  - Uses fine-grain locks and a concurrent priority queue

```
Graph (E,V,w)
done[1..TNum] = [false,...,false]
D[1..|V|] = [ $\infty$ ,..., $\infty$ ]
Element[1..|V|] Offer =
    [null,...,null]
Lock [1.. |V|] DLock
Lock [1.. |V|] OfferLock
```

```
relax(v,vd)
    lock(OfferLock[v])
    if (vd < D[v])
        vo = Offer[v]
        if (vo == null)
            Offer[v] = insert(v,vd)
        else
            if (vd < vo.key)
                Offer[v] = insert(v,vd)
    unlock(OfferLock[v])
```

```
parallelDijkstra()
    while (!done[tid])
        o = extractMin()
        if (o  $\neq$  null)
            u = o.data
            d = o.key
            lock(DLock[u])
            if (dist < D[u])
                D[u] = d
                explore = true
            else
                explore = false
            unlock(DLock[u])
            if (explore)
                foreach ((u,v)  $\in$  E)
                    vd = d + w(u,v)
                    relax(v,vd)
        else
            done[tid] = true
            i = 0
            while (done[i] and i < TNum)
                i = i + 1
            if (i == TNUM)
                return
            done[tid] = false
```

# Priority Queue using Multi-Dimension Lists

## Key Conversion

- key  $\rightarrow (\text{key})_{\lceil \sqrt[D]{N} \rceil} \rightarrow [k_0, k_1, \dots, k_{D-1}]$

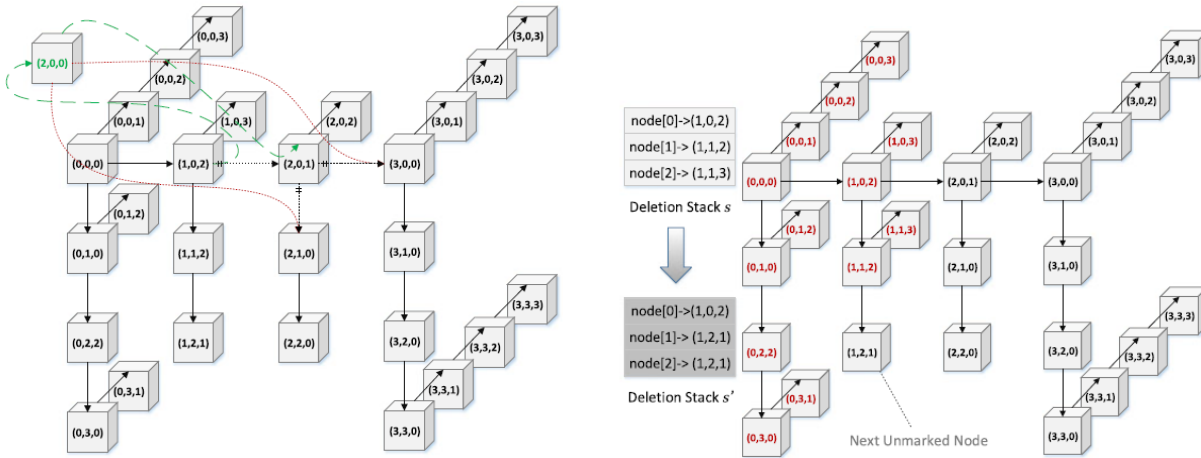
## Basic Rule

- Dimension range:  $[0, D)$ ; each node has up to D children
- The coordinates of a child is greater than or equal to the coordinate of the parent for all dimensions

## Data Structures

- Adesc (Adoption Description)
  - parent node of adoption
  - child slot to be adopted (dp, dc)
- Node
  - key, value
  - key coordinates
  - child[D]
  - adesc
- Stack (deletion stack to trace the last known deleted node)
- Priority Queue
  - D, N
  - head (dummy head)
  - stack


# Insert



- Build the new node
- LocatePred  $\rightarrow$  predNode, childNode, dp, dc
- FinishInserting  $\rightarrow$  child adoption for predNode, childNode
- Fill in the new node(CAS)
- Rewind the deletion stack
- Retry if 1) another thread inserted a child; 2) the child slot has been marked as invalid

## Rewind the stack

- Synchronize info between insert and deleteMin
- The old stack points to a position beyond the new node



## DeleteMin

- Backtrack the deletion stack to find the next minimum stack
- Retry if the found node has been marked as deleted


# Implementation Extensions

- Duplicated priorities
  - Reserve 13/32 bits in key space for unique identifiers
  - Identifiers are assigned using an atomic counter
- DeleteMin return values
  - SSSP requires the priority and value of the min entry
  - SSSP requires to know when priority queue is empty
- Support non-reference values
  - Reference implementation assumes stored values are referenced as pointers
- Exclusion of physical deletion
  - The reference paper indicated they use only logical deletion for benchmarks
- Pseudocode bugs
  - There were several instances of wrong or missing pseudocode from both the MDList and Parallel Dijkstra algorithms


# Results

## Microbenchmarks


GHC 100% Insert




PSC 100% Insert




GHC 100% Delete




PSC 100% Delete



GHC Mixed




PSC Mixed

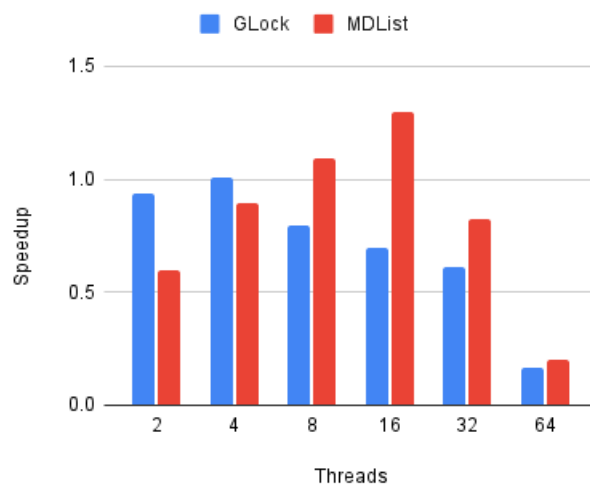


# SSSP Benchmark


## GHC 1024 Nodes




## PSC 1024 Nodes




## GHC 4096 Nodes




## PSC 4096 Nodes



## GHC 8191 Nodes



## PSC 8191 Nodes



# Discussion and Analysis

## Microbenchmarks

- Coarse grain lock performs poorly with more threads
  - Due to more lock contention
- MDList outperforms coarse grain for  $\geq 4$  threads
  - Insert performs poorly  $< 8$  threads due to overhead
  - Insert performs well relatively for 8 - 32 threads
  - Insert performs poorly  $> 32$  due to context switching
  - Insert bottlenecked by allocation (via VTune results)
  - Delete performance constant
    - Due to inherent sequential nature of DeleteMin
  - Mixed performance inconclusive since coarse grain implements physical deletion while MDList does not

## SSSP Benchmark

- Coarse grain lock performs poorly with more threads
  - Due to more lock contention
  - Speedup mostly attributed to parallelized SSSP
- MDList outperforms coarse grain for  $\geq 4$  threads on GHC and  $\geq 8$  threads on PSC
  - Performs poorly for low thread counts due to overhead
  - Scales well to up to 32 threads on PSC
  - DeleteMin bottlenecked by CAS retry (via diagnostics)
- Superlinear speedup on GHC due to lower branch miss rate and cache miss rate based on prof result

# Further Work

- Concurrent memory allocation
  - Experiment with other allocators instead of glibc [3]
- Comparisons with alternative priority queue implementations
  - Intel's TBBPQ [4]
  - Fine Grained Locking [5]
- Physical Deletion
  - Implement Purge
- Further Optimizations of MDList implementation

# Conclusion

- Extended MDList to solve Parallelized SSSP
- Observed performance improvement over coarse grained locking on SSSP benchmarks and microbenchmarks
- Limitations of concurrent priority queues limits its application in highly parallelized algorithms [6]

# References

- [1] Zhang, D. and Dechev, D. A lock-free priority queue design based on multi-dimensional linked lists. *IEEE Transactions on Parallel and Distributed Systems*, 27(3):613–626, 2016. doi: 10.1109/TPDS.2015.2419651.
- [2] Tamir, O., Morrison, A., and Rinetzky, N. A heap-based concurrent priority queue with mutable priorities for faster parallel algorithms. In *OPODIS*, 2015.
- [3] Evans, J. A scalable concurrent malloc (3) implementation for freebsd. In *Proc. of the bsdcan conference*, Ottawa, Canada, 2006.
- [4] Shavit, N. and Lotan, I. Skiplist-based concurrent priority queues. In *Proceedings 14th International Parallel and Distributed Processing Symposium. IPDPS 2000*, pp. 263–268. IEEE, 2000.
- [5] Hunt, G. C., Michael, M. M., Parthasarathy, S., and Scott, M. L. An efficient algorithm for concurrent priority queue heaps. *Information Processing Letters*, 60(3):151–157, 1996.
- [6] Srinivasan, S., Riazzi, S., Norris, B., Das, S. K., and Bhowmick, S. A shared-memory parallel algorithm for updating single-source shortest paths in large dynamic networks. In *2018 IEEE 25th International Conference on High Performance Computing (HiPC)*, pp. 245–254, 2018. doi: 10.1109/HiPC.2018.00035.